


DEPARTMENT OF INFORMATICS


TECHNISCHE UNIVERSITÄT MÜNCHEN

Bachelor's Thesis in Informatics

**Training a Spiking Neural Network using
R-STDP to perform Autonomous Target
Tracking on a Snake Car Robot**

René Romen





DEPARTMENT OF INFORMATICS


TECHNISCHE UNIVERSITÄT MÜNCHEN

Bachelor's Thesis in Informatics

Training a Spiking Neural Network using R-STDP to
perform Autonomous Target Tracking on a Snake Car
Robot

Training von Spiking Neural Networks für die
Automatische Zielverfolgung auf einem Schlangen
Roboter mit R-STDP

Author: René Romen
Supervisor: Prof. Dr.-Ing. habil. Alois Knoll
Advisors: Zhenshan Bing, M.Sc.
Date: August 6, 2018



I confirm that this master's thesis is my own work and I have documented all sources and material used.

München, 6. August 2018

René Romen

Abstract

Contents

Abstract	vii
List of Abbreviations	xi
1. Introduction	1
1.1. Background	1
1.2. Problem Statement	1
2. Technology Overview	3
3. Implementation ...	5
4. Implementation of Direct GPU-FPGA Communication	7
4.1. GPUDirect RDMA Overview	7
4.2. Extension of the Altera PCIe Driver	8
4.2.1. Basic Version	9
4.2.2. Optimizations	11
4.3. GPUDirect RDMA for OpenCL	13
4.3.1. Reverse Engineering the NVIDIA Driver Communication	13
4.3.2. Extension of the NVIDIA Kernel Module	16
4.4. User Space Invocation	17
5. Implementation of Concurrent Indirect GPU-FPGA Communication	21
6. Evaluation	25
6.1. Hardware Configuration	25
6.2. Effects of RDMA Optimizations	25
6.3. Parameter Choice for Concurrent Indirect Transfer	27
6.4. Method Comparison	27
6.5. Comparison with Previous Work	28
7. Conclusions and Future Work	31
Appendix	35
A. Example RDMA application	35
B. Setup Instructions	37

Bibliography	39
---------------------	-----------

List of Abbreviations

SNN Spiking Neural Network

1. Introduction

1.1. Background

1.2. Problem Statement

This thesis is structured as follows:

- Chapter 2 provides ...

2. Technology Overview

3. Implementation of Direct GPU-FPGA Communication

Listing 3.1: example

```
code = example
```


4. Evaluation

5. Conclusions and Future Work

Appendix

A. Example RDMA application

Bibliography

- [1] Altera Corporation. FPGA Architecture - White Paper. https://www.altera.com/en_US/pdfs/literature/wp/wp-01003.pdf, 2006. Retrieved on 2015.03.13.
- [2] Altera Corporation. Embedded Peripherals IP User Guide. https://www.altera.com/en_US/pdfs/literature/ug/ug_embedded_ip.pdf, 2011. Retrieved on 2015.03.24.
- [3] Altera Corporation. Altera SDK for OpenCL (version 13.1). <http://dl.altera.com/openc1/>, 2013. Retrieved on 2015.03.28.
- [4] Altera Corporation. Implementing FPGA Design with the OpenCL Standard. <http://www.altera.com/literature/wp/wp-01173-openc1.pdf>, 2013. Retrieved on 2015.03.13.
- [5] Altera Corporation. Embedded Peripheral IP User Guide, 2014. Chapter 22: Altera Modular Scatter-Gather DMA.
- [6] Altera Corporation. Stratix V Avalon-MM Interface for PCIe Solutions User Guide, 2014.
- [7] Altera Corporation. Avalon Interface Specifications. https://www.altera.com/en_US/pdfs/literature/manual/mnl_avalon_spec.pdf, 2015. Retrieved on 2015.04.05.
- [8] Ray Bittner. Speedy Bus Mastering PCI Express. *22nd International Conference on Field Programmable Logic and Applications (FPL 2012)*, August 2012.
- [9] Ray Bittner and Erik Ruf. Direct GPU/FPGA Communication Via PCI Express. *1st International Workshop on Unconventional Cluster Architectures and Applications (UCAA 2012)*, September 2012.
- [10] Nick Black. Libcudest. <http://nick-black.com/dankwiki/index.php/Libcudest>. Retrieved on 2015.03.18.
- [11] Ravi Budruk. *PCI express system architecture*. Addison-Wesley, 2003.
- [12] NVIDIA Corporation. NVIDIA GPUDirect. <https://developer.nvidia.com/gpudirect>, 2010. Retrieved on 2015.04.02.
- [13] NVIDIA Corporation. NVIDIA's Next Generation CUDA(TM) Compute Architecture: Kepler(TM) GK110. <http://www.nvidia.com/content/PDF/kepler/NVIDIA-Kepler-GK110-Architecture-Whitepaper.pdf>, 2012. Retrieved on 2015.02.25.

- [14] NVIDIA Corporation. Developing a Linux Kernel Module using GPUDirect RDMA. <http://docs.nvidia.com/cuda/gpudirect-rdma/index.html>, 2014. Retrieved on 2015.03.11.
- [15] NVIDIA Corporation. CUDA C Programming Guide [version 7.0]. http://docs.nvidia.com/cuda/pdf/CUDA_C_Programming_Guide.pdf, 2015. Retrieved on 2015.04.03.
- [16] Umer Farooq, Zied Marrakchi, and Habib Mehrez. *Tree-based Heterogeneous FPGA Architectures - Application Specific Exploration and Optimization*. Springer, Berlin, Heidelberg, 2012.
- [17] Scott Hauck. *Reconfigurable computing the theory and practice of FPGA-based computation*. Morgan Kaufmann, Amsterdam Boston, 2008.
- [18] Hewlett-Packard. Hewlett-Packard Quadro K600 Specifications. <http://www8.hp.com/h20195/v2/GetDocument.aspx?docname=c04128134>, 2013. Retrieved on 2015.04.07.
- [19] Ra Inta, David J. Bowman, and Susan M. Scott. The “Chimera”: An Off-The-Shelf CPU/GPGPU/FPGA Hybrid Computing Platform, 2012. International Journal of Reconfigurable Computing, Volume 2012.
- [20] Khronos Group. OpenCL Extension #5: Installable Client Driver (ICD) Loader. https://www.khronos.org/registry/cl/extensions/khr/cl_khr_icd.txt, 2010. Retrieved on 2015.03.15.
- [21] Khronos Group. OpenCL 1.2 Installable Client Driver Loader [Source]. <https://www.khronos.org/registry/cl/specs/opencl-icd-1.2.11.0.tgz>, 2012. Retrieved on 2015.04.01.
- [22] Khronos Group. The OpenCL Specification Version 1.2. <https://www.khronos.org/registry/cl/specs/opencl-1.2.pdf>, 2012. Retrieved on 2015.03.15.
- [23] Linux man-pages project. dlopen(3) Linux man page, 2000.
- [24] Linux man-pages project. ioctl(2) Linux man page, 2000.
- [25] Nikhil Madduri. Hardware Accelerated Particle Filter for Lane Detection and Tracking in OpenCL. Master’s thesis, Technische Universität München, 2014.
- [26] Nallatech Inc. Nallatech 385-D5 Specifications. http://www.nallatech.com/wp-content/uploads/pcie_385pb_v1_21.pdf, 2014. Retrieved on 2015.04.07.
- [27] NVIDIA Corporation. NVIDIA OpenCL Best Practices Guide. http://www.nvidia.com/content/cudazone/CUDABrowser/downloads/papers/NVIDIA_OpenCL_BestPracticesGuide.pdf, 2009. Retrieved on 2015.03.21.
- [28] NVIDIA Corporation. LINUX X64 (AMD64/EM64T) DISPLAY DRIVER (Version 346.47) [kernel/nv.c]. http://us.download.nvidia.com/XFree86/Linux-x86_64/346.47/NVIDIA-Linux-x86_64-346.47.run, 2015. Retrieved on 2015.03.19.

- [29] Alessandro Rubini, Jonathan Corbet, and Greg Kroah-Hartman. *Linux Device Drivers, 3rd Edition*. O'Reilly Media, 2005.
- [30] Shinpei Kato and Michael McThrow and Carlos Maltzahn and Scott Brandt. Gdev: First-Class GPU Resource Management in the Operating System. Tech. Rep., 2012, uSENIX Annual Technical Conference (USENIX ATC'12).
- [31] David Susanto. Parallelism for Computationally Intensive Algorithm with GPU/FPGA Interleaving. Master's thesis, Technische Universität München, 2014.
- [32] Andrew S Tanenbaum. *Modern Operating Systems*. Pearson, 2 edition, 2007.
- [33] Yann Thoma, Alberto Dassatti, and Daniel Molla. FPGA²: An open source framework for FPGA-GPU PCIe communication. In *2012 International Conference on Reconfigurable Computing and FPGAs, ReConFig 2013, Cancun, Mexico, December 9-11, 2013*, pages 1–6, 2013.
- [34] Various Authors. Nouveau: Accelerated Open Source driver for nVidia cards. <http://nouveau.freedesktop.org/wiki/>, 2012. Retrieved on 2015.04.10.
- [35] Various Authors. Nouveau: Feature Matrix. <http://nouveau.freedesktop.org/wiki/FeatureMatrix/>, 2014. Retrieved on 2015.04.11.
- [36] Nicholas Wilt. *The CUDA handbook a comprehensive guide to GPU programming*. Addison-Wesley, Upper Saddle River, NJ, 2013.